

Captive Portal / Hotspot Design

Internal Architecture Spec: Hotspot Profile Deployment System [↗](#)

Objective [↗](#)

To provide clients with a backend tool that automates the configuration and deployment of secure, traceable captive portal environments. The system integrates with CoovaChilli, FreeRADIUS, and optionally Squid, based on client-defined Hotspot Profiles. Aruba Access Points serve as the external SSID broadcasters and are outside our management scope.

System Components [↗](#)

1. Aruba Access Point (External) [↗](#)

- Broadcasts Wi-Fi SSID (e.g., "Guest-WiFi")
- Places connected clients in a specified VLAN (e.g., 182)
- Handles DNS/DHCP unless otherwise specified

2. Hotspot Profile Manager (Core Backend) [↗](#)

- Allows definition of network-facing Hotspot Profiles
- Generates and applies configurations for:
 - CoovaChilli
 - FreeRADIUS
 - Squid (optional)
- Provides secure defaults and constraints to ensure valid setups

3. CoovaChilli [↗](#)

- Captive portal controller
- Manages client redirection, IP routing, session initiation
- Communicates with FreeRADIUS for authentication/authorization
- Optional integration with Squid for HTTP proxying

4. FreeRADIUS [↗](#)

- Authenticates users (voucher, username/password, MAC)
- Stores accounting logs for tracking session data and identity
- Enables full traceability of network access

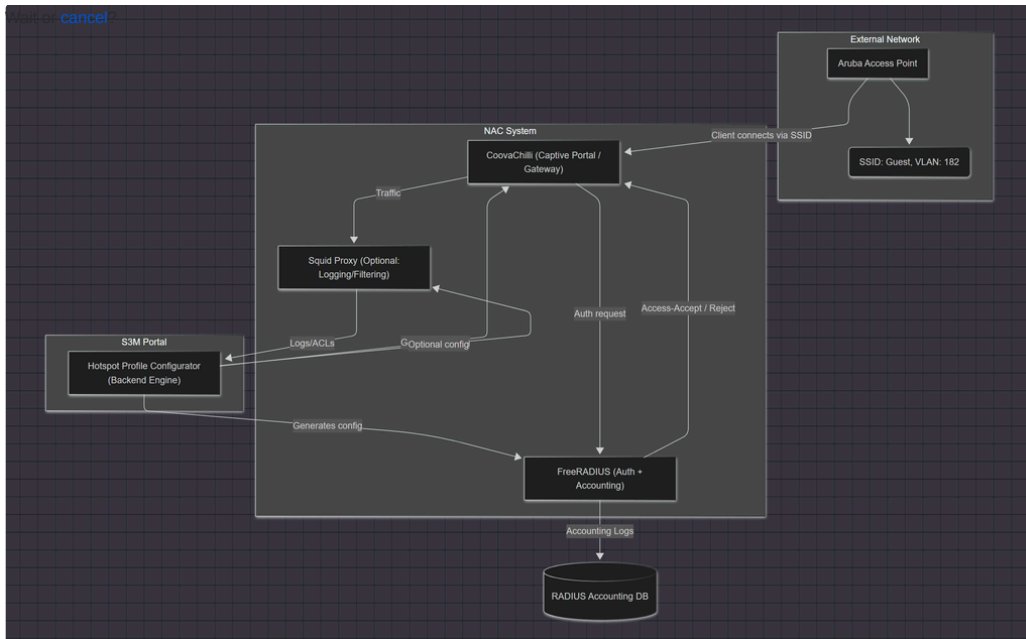
5. Squid (Optional) [↗](#)

- HTTP/HTTPS proxy for traffic logging and access control
- Enforced via transparent iptables NAT rules from CoovaChilli
- Adds an additional layer of monitoring or filtering if required

Architecture Diagram [🔗](#)

App is not responding. Wait or [cancel?](#)

App is not responding. [cancel](#)



Data and Traffic Flow [🔗](#)

- Client connects to Aruba AP**
 - Receives IP via DHCP (from Aruba or firewall)
 - Placed in Guest VLAN (e.g., 182)
- CoovaChilli intercepts Traffic**
 - Redirects HTTP to captive portal (hosted locally or pre-generated)
 - Handles IP/MAC mapping and tracks session
- User Authenticates**
 - CoovaChilli sends RADIUS Access-Request
 - FreeRADIUS replies (Accept/Reject)
 - CoovaChilli enables forwarding based on response
- Optional: Squid Proxy**
 - HTTP/HTTPS traffic routed to Squid via NAT
 - Squid logs domain/IP and maps to client identity
- RADIUS Accounting**
 - Start/Stop packets track session time, data usage
 - Stored in backend DB for traceability and reporting

System Overview [🔗](#)

Building a **backend system** that allows clients to define **Hotspot Profiles** through a database. Each profile describes how to configure a local instance of:

- CoovaChilli:** the captive portal gateway
- FreeRADIUS:** the authentication/accounting backend

- **Squid** (optional): HTTP/HTTPS logging proxy

This profile-driven design dynamically generates and deploys working configurations **locally on the same server**.

Component Responsibilities [↗](#)

Component	Source of Config	Deployment Target	Notes
CoovaChilli	DB → generate_coova_config(profile)	/etc/chilli/defaults	Handles portal redirect and DHCP-like behavior
FreeRADIUS	DB → generate_radius_config(profile)	/etc/raddb/clients.conf, /etc/raddb/users, or SQL ⚠	Handles authentication and session accounting
Squid (optional)	DB → generate_squid_config(profile)	/etc/squid/squid.conf f	Logs and filters post-auth HTTP/HTTPS

Hotspot Profile Database Design (Draft Schema) [↗](#)

```
1 CREATE TABLE hotspot_profiles (  
2     id SERIAL PRIMARY KEY,  
3     name TEXT UNIQUE,  
4     vlan_id INTEGER NOT NULL,  
5     interface TEXT NOT NULL,  
6     network_cidr TEXT NOT NULL,           -- e.g., '192.168.182.0/24'  
7     uam_server TEXT NOT NULL,             -- e.g., 'http://192.168.182.1:3990'  
8     uam_secret TEXT NOT NULL,  
9     uam_portal_type TEXT NOT NULL CHECK (uam_portal_type IN ('radius_login', 'voucher')),  
10    auth_method TEXT NOT NULL CHECK (auth_method IN ('radius', 'voucher')),  
11    radius_ip TEXT NOT NULL,  
12    radius_secret TEXT NOT NULL,  
13    proxy_enabled BOOLEAN DEFAULT FALSE,  
14    proxy_port INTEGER DEFAULT 3128,  
15    allowed_domains TEXT[],               -- JSON or array, e.g. ['example.com']  
16    created_at TIMESTAMP DEFAULT now()  
17 );
```

Note: Secure fields like `radius_secret` and `uam_secret` should be **encrypted** or at least hashed-at-rest if planned to be shown in the frontend.

Field-by-Field Explanation [↗](#)

Field	Why It's Important	Where It's Used	Example Data
profile_name	Used to identify the profile	Naming configs, NAS-ID in RADIUS	cafe_latte_guest

<code>vlan_id</code>	Matches Aruba's Guest VLAN	Used for tagging/logging only (not needed in config)	182
<code>interface</code>	Physical NIC or VLAN interface that CoovaChilli will bind to	<code>HS_INTERFACE</code> in Coova	eth0
<code>network_cidr</code>	Guest IP pool for clients	<code>HS_NETWORK</code> , <code>HS_NETMASK</code> in Coova	192.168.182.0/24
<code>uam_server</code>	IP/URL of captive portal	Used in CoovaChilli's <code>HS_UAMSERVER</code> , redirect logic	http://192.168.182.1:3990
<code>uam_https</code>	If there is a valid certificate to be used	<code>HS_UAMHTTPS</code> for the captive portal in Coova	on off
<code>uam_secret</code>	Shared secret for UAM portal integrity	<code>HS_UAMSECRET</code> in Coova, and used in captive page submission	abcd1234
<code>auth_method</code>	Whether FreeRADIUS handles user/pass or vouchers	Chooses logic for FreeRADIUS <code>users</code> file or SQL	radius voucher
<code>radius_ip</code>	Where to send RADIUS requests (local, but still explicitly set)	<code>HS_RADIUS</code> in CoovaChilli config	10.34.10.5
<code>radius_secret</code>	Shared secret between Coova and RADIUS	Used in both CoovaChilli and FreeRADIUS <code>clients.conf</code> (Creating Access Device)	testing123
<code>proxy_enabled</code>	Whether to activate Squid proxying	Triggers Squid setup and iptables redirect	true false
<code>proxy_port</code>	Port for Squid to listen on	<code>3128</code> by default, injected into Squid config	3128
<code>allowed_domains</code>	Domains allowed before authentication (e.g., CDN, Google login)	<code>HS_UAMALLOWED</code> in CoovaChilli	['api.example.com', 'static.examplecdn.com']

Goal:

When a new hotspot profile is added to the database, the system should:

1. Automatically create and deploy a new **CoovaChilli instance**
2. FreeRADIUS must respond to **Access-Request** and **Accounting-Request** messages from CoovaChilli.
3. **Enhances visibility and control** over user browsing after login with Squid if requested

Workflow Overview [↗](#)

```
1 1. Admin creates hotspot profile → stored in DB
2 2. Backend validates input
3 3. Each component generator reads the DB and generates its config:
4   - CoovaChilli: /etc/chilli/defaults
5   - FreeRADIUS: /etc/raddb/clients.conf, /etc/raddb/users or DB
6   - Squid (optional): /etc/squid/squid.conf
7 4. System restarts services:
8   - systemctl restart chilli
9   - systemctl stop freeradius
10  - systemctl start freeradius
11  - systemctl restart squid (if enabled)
```

Big Picture Workflow (Step-by-Step) [↗](#)

1. A New Hotspot Profile is Created [↗](#)

- A user or admin adds a new hotspot profile via CLI, API, or directly in the DB.
- Example: a profile called `cafe_latte` with interface `eth0.100` (VLAN 100).

2. ConfigController Detects a New Profile [↗](#)

- A REST API handler
- The controller receives `profile_id` (or the full profile dict)

3. ConfigController Starts Deployment Pipeline [↗](#)

- It calls:
 - ```
1 deploy_profile(profile_id)
```
- Which internally triggers:
  - ```
1 - deploy_coova_config(profile)
2 - deploy_radius_config(profile)
3 - deploy_squid_config(profile)
```

4. CoovaChilli Configuration Phase [↗](#)

`deploy_coova_config(profile)` does the following:

Sub-Step	Description
Parse the profile	Reads info like interface name, UAM server URL, IP range, RADIUS IP
Generate a config file	Uses the profile to fill out a template like <code>cafe_latte.conf</code>

Link it for systemd	Symlink this config to <code>/etc/chilli/cafe_latte.conf</code>
Register it as a systemd instance	Coova will run as <code>chilli@cafe_latte.service</code>
Start the service	<code>systemctl start chilli@cafe_latte</code> creates a running captive portal gateway

 It is crucial that **CoovaChilli can run multiple instances simultaneously**, as long as each instance:

- Runs as a **separate process**
- Has its **own config file**
- Binds to a **unique physical or virtual interface**
- Uses a **different UAM listen IP and tun interface**
- Avoids port conflicts

This approach is **used in production environments**, especially for multi-tenant networks, public hotspots, and ISPs.

Verified Conditions for Multiple CoovaChilli Instances

Here's what you must ensure for this to actually work:

Config Element	Must Be
<code>HS_INTERFACE</code>	Unique per instance (e.g., <code>eth0.100</code> , <code>eth0.200</code>)
<code>HS_NETWORK</code> / <code>HS_NETMASK</code>	Non-overlapping IP pools per instance
<code>HS_UAMLISTEN</code>	Unique IP per instance (e.g., <code>192.168.182.1</code> , <code>192.168.183.1</code>)
<code>tunX</code> interface	Automatically created (e.g., <code>tun0</code> , <code>tun1</code> , etc.)
<code>iptables</code> rules	Written separately per instance (ideally per <code>tunX</code>)
<code>systemd</code>	Use template units: <code>chilli@profile_name.service</code>

File-Level Isolation

Each instance should have:

- Its own config: `/etc/chilli/cafe_latte.conf` , `/etc/chilli/techfair.conf`
- Its own log files (optional, if configured)
- Its own IP range, VLAN interface, and redirect portal

Things to Watch Out For

- **UAM ports** can conflict (default `3990`) → but only if using same IP/UAM listener
- You **must manage iptables/NAT rules separately** for each instance's `tunX` device
- Each interface (e.g., `eth0.100`) must be routable to your guests
- Ensure **routing/NAT/firewall rules** are written with interface awareness

5. Recognize CoovaChilli as a valid NAS [↗](#)

- Coova sends requests as a RADIUS client (NAS)
- FreeRADIUS must trust it as a NAS device, such as in the NAS table in the database

6. Handle Authentication Requests [↗](#)

Coova sends:

```
1 Access-Request
2   User-Name = "guest123"
3   User-Password = "secret"
4   Calling-Station-Id = "MAC_ADDRESS"
5   NAS-Identifier = "cafe_latte"
```

FreeRADIUS must proceed to guest services:

- Look up this user
- Validate credentials (username/password or voucher)
- Return an **Access-Accept** (or **Access-Reject**) with **optional attributes**

Depending on your profile type:

- If it's **voucher-based**: FreeRADIUS should check that the voucher exists, is unused/valid, then mark it as used
- If it's **username/password**: validate against guest user table, LDAP, or flat file

7. Return Proper Reply Attributes [↗](#)

You can control the session using reply attributes:

Attribute	What It Does
Session-Timeout	Disconnect the user after X seconds
Idle-Timeout	Disconnect the user after X seconds of inactivity
ChilliSpot-Max-Total-Octets	Data cap in bytes
WISPr-Bandwidth-Max-Down / Up	Rate limits in bps

Example Access-Accept:

```
1 Access-Accept
2   Session-Timeout = 1800
3   ChilliSpot-Max-Total-Octets = 50000000
```

8. Handle Accounting Packets [↗](#)

Coova will send:

```
1 Accounting-Start
2 Accounting-Interim-Update
3 Accounting-Stop
```

FreeRADIUS must:

- Log these (in SQL or files)
- Track session start/stop time, IP, MAC, data used
- Tie this data back to the profile (e.g., via NAS-Identifier or Calling-Station-Id)

This is **how traffic is traced back to a guest** — these logs are mandatory!

If you use a SQL-based accounting backend (e.g., `rlm_sql`), check that your schema includes:

- username
- framed_ip
- calling_station_id (MAC)
- acct_session_id
- acct_input_octets / output_octets

9. Configure Squid [↗](#)

In `/etc/squid/squid.conf`:

```
1 # Listen on default proxy port
2 http_port 3128 transparent
3
4 # Define ACLs
5 acl allowed_localnet src 192.168.182.0/24 # Change to match Coova subnet
6 http_access allow allowed_localnet
7
8 # Logging format (optional but recommended)
9 logformat custom_with_ip %ts.%03tu %>a %un %rm %ru %Hs %<st
10 access_log /var/log/squid/access.log custom_with_ip
```

- `http_port 3128 transparent` enables intercept mode
- `%>a` logs client IP, which is **critical for attribution**
- You can extend logging to include usernames if available

10. Enable NAT Redirection via iptables [↗](#)

After a user logs in and is allowed to surf:

```
1 iptables -t nat -A PREROUTING -i tun0 -p tcp --dport 80 -j REDIRECT --to-port 3128
```

This command captures all HTTP traffic from clients going through `tun0` (Coova's internal interface) and sends it to Squid.

You can add similar rules for port 443 (HTTPS), but **it requires SSL bumping** — advanced and potentially privacy-sensitive.

11. Access Logs [↗](#)

Squid will log traffic like:

```
1 1681328475.123 56 192.168.182.101 TCP_MISS/200 1652 GET http://example.com/ - DIRECT/93.184.216.34 text/html
```

You can parse logs to match:

- IP → `radacct.framedipaddress`

- MAC → via CoovaChilli session table
- Username → via RADIUS session or accounting

What Is a Voucher? [🔗](#)

A **voucher** is simply a **username** (and optionally a password) that:

- Has a **limited lifetime** (time-based or expiration date)
- Is usually **pre-generated**
- Can be **single-use or multi-use**
- Is stored in a **database or flat file**
- Can return **custom RADIUS reply attributes** (like time or data caps)

Voucher System Architecture (Simplified) [🔗](#)

```

1  [CoovaChilli Portal Page]
2      → User enters voucher code
3      ↓
4  [FreeRADIUS]
5      → Checks voucher code against database or file
6      → If valid:
7          → Returns Access-Accept + Session-Timeout + bandwidth limits
8          → Marks voucher as used (if one-time)
9      → If invalid/expired:
10         → Returns Access-Reject

```

Voucher Lifecycle Flow [🔗](#)

1. **Create vouchers** (CLI or portal)
 - E.g., generate 100 codes, each valid for 1 hour, expires in 30 days
2. **Store them in a database** (SQL) or file
 - Fields: code, used, created_at, expires_at, session_time, octet_limit
3. **User enters code in captive portal**
 - Form submits username (voucher code) to CoovaChilli
 - CoovaChilli passes it as `User-Name` to RADIUS
4. **FreeRADIUS authenticates**
 - Finds the voucher
 - Verifies it is not used or expired
 - Replies with Access-Accept and session control
5. **FreeRADIUS marks it as used**
 - For one-time vouchers, `used = true` or deletes it
6. **RADIUS accounting tracks usage**
 - Username, IP, MAC, duration, bandwidth

SQL-Based Voucher System [🔗](#)

1. **Create a SQL table**

```

1 CREATE TABLE vouchers (
2     code VARCHAR(64) PRIMARY KEY,
3     created_at TIMESTAMP DEFAULT now(),
4     expires_at TIMESTAMP,
5     session_time INTEGER DEFAULT 1800, -- in seconds
6     octet_limit BIGINT DEFAULT NULL, -- optional
7     is_used BOOLEAN DEFAULT false
8 );

```

2. Enable SQL module in FreeRADIUS

- Use `rlm_sql` and define queries in `mods-config/sql/authorize-check.sql`

3. Authorize with SQL

- Add a policy in `sites-enabled/default` → `authorize`:

```

1 authorize {
2     ...
3     sql
4     if (!ok) {
5         reject
6     }
7 }

```

4. Control usage in `authorize`

```

1 if ("%{sql:SELECT is_used FROM vouchers WHERE code='%{User-Name}%'}" == "t") {
2     reject
3 }

```

5. Add reply attributes dynamically

```

1 update reply {
2     Session-Timeout := "%{sql:SELECT session_time FROM vouchers WHERE code='%{User-Name}%'}"
3 }

```

6. Mark voucher as used (Post-Auth or Accounting-Start)

- Hook into `post-auth` or `sql` module's `post-auth-query`:

```

1 UPDATE vouchers SET is_used=true WHERE code='%{User-Name}';

```

Voucher Customization Examples [🔗](#)

Field	Purpose
<code>session_time</code>	Auto-expire after X seconds
<code>expires_at</code>	Prevent login after a certain date
<code>is_used</code>	Single-use control
<code>octet_limit</code>	Cap on total data usage
<code>bandwidth_up/down</code>	Per-user bandwidth limits
<code>group</code>	Group vouchers by location or SSID/profile

Health Check [🔗](#)

Purpose: [🔗](#)

To detect if a **CoovaChilli instance (or system component)** is:

- Misconfigured
- Crashed
- Hung
- Not serving clients properly

Health Check (Per Instance): [🔗](#)

Check	How	Why
CoovaChilli process running	<code>systemctl is-active chilli@profile</code>	Detects crash or failure to start
tunX interface exists	<code>ip link show tun0 (or tun1...)</code>	Verifies tunnel is up
UAM server reachable	<code>curl -k https://portal.example.com</code>	Detects portal web server failure
Session check	<code>chilli_query list</code> has output	Confirms users are able to connect

Sample Concept [🔗](#)

```
1 check_chilli_health() {
2   profile=$1
3   systemctl is-active chilli@$profile || echo "$profile is down"
4   ip link show tun0 >/dev/null || echo "Tunnel not up"
5   curl -k --max-time 3 https://portal.example.com/login || echo "Portal unreachable"
6 }
```

Sample Code Structure [🔗](#)

```
1 hotspot_manager/
2 |
3 |─ main.py                # Entrypoint (CLI or REST)
4 |
5 |─ config_generators/     # Handles config generation for each tool
6 |   |─ __init__.py
7 |   |─ coova.py           # CoovaChilli generator, deploy, delete
8 |   |─ radius.py         # FreeRADIUS generator, deploy, delete
9 |   |─ squid.py          # Squid generator, deploy, delete
10 |
11 |─ controllers/          # Central orchestration logic
12 |   |─ __init__.py
13 |   |─ config_controller.py # deploy_profile(), update_profile(), delete_profile()
14 |
15 |─ templates/            # Jinja2 templates for config generation
```

```

16 |   └─ coova_defaults.j2
17 |   └─ radius_clients.j2
18 |   └─ squid_conf.j2
19 |
20 └─ utils/                                # Utilities
21 |   └─ __init__.py
22 |   └─ db.py                            # DB functions: get_profile_by_id(), get_all_profiles()
23 |   └─ netutils.py                      # parse_cidr(), validate_ip(), calculate_netmask()
24 |   └─ templates.py                    # Jinja2 wrapper: render_template(template_name, context)
25 |
26 └─ migrations/                          # SQL scripts for creating and upgrading the DB schema
27 |   └─ create_hotspot_profiles.sql
28 |
29 └─ logs/                                # Config deployment logs
30   └─ deploy.log

```

Key Modules [🔗](#)

config_generators/coova.py [🔗](#)

- generate_coova_config(profile)
- deploy_coova_config(profile)
- delete_coova_config(profile_name)

controllers/config_controller.py [🔗](#)

- deploy_profile(profile_id)
- update_profile(profile_id)
- delete_profile(profile_id)

utils/db.py [🔗](#)

- get_profile_by_id(profile_id)
- insert_profile(data)
- update_profile(profile_id, data)
- delete_profile(profile_id)

utils/netutils.py [🔗](#)

- parse_cidr('192.168.182.0/24') → ('192.168.182.0', '255.255.255.0')

Lifecycle Handlers [🔗](#)

Action	Trigger	Operation
Add Profile	POST /profiles or create_profile()	Validate + generate config + deploy
Edit Profile	PUT /profiles/:id or update_profile()	Regenerate config + restart affected service(s)
Delete Profile	DELETE /profiles/:id or delete_profile()	Remove config file + restart/clean affected services

