

# FreeRADIUS & rlm\_python

## What is rlm\_python ? [🔗](#)

It's a FreeRADIUS module that lets you:

- Intercept `authorize`, `authenticate`, `post-auth`, `accounting`, etc.
- Execute arbitrary Python logic during these phases
- Read request attributes and modify reply or control attributes
- Dynamically pull data from a DB or API

### 1. Install rlm\_python [🔗](#)

On Debian/Ubuntu:

```
1 sudo apt install freeradius-python3
```

### 2. Enable the Module [🔗](#)

```
1 cd /etc/freeradius/3.0/mods-enabled/  
2 ln -s ../mods-available/python python
```

### 3. Configure the Module [🔗](#)

Edit:

```
/etc/freeradius/3.0/mods-available/python
```

Set:

```
1 python {  
2     module = policy_engine  
3     python_path = ${modconfdir}/${!.:name}/site-packages  
4 }
```

Then create the directory:

```
1 mkdir -p /etc/freeradius/3.0/mods-config/python/site-packages/
```

### 4. Python File (policy\_engine.py) [🔗](#)

Create:

```
/etc/freeradius/3.0/mods-config/python/site-packages/policy_engine.py
```

Example skeleton:

```
1 def authorize(p):  
2     username = p["User-Name"]  
3     nas_ip = p.get("NAS-IP-Address", "unknown")  
4  
5     # Load policies from DB (you can use sqlalchemy or pycpg2 here)  
6     policy = match_policy(username, nas_ip, p)  
7  
8     if not policy:
```

```

9         return 1 # Reject
10
11     backend = get_backend(policy['auth_backend'])
12
13     success, user_data = backend.authenticate(username, p["User-Password"])
14     if not success:
15         return 1 # Reject
16
17     if policy.get("check_blacklist", False):
18         if is_blacklisted(username, user_data["mac"]):
19             return 1 # Reject
20
21     # VLAN reply attribute example
22     p["reply:Tunnel-Type"] = "VLAN"
23     p["reply:Tunnel-Medium-Type"] = "IEEE-802"
24     p["reply:Tunnel-Private-Group-Id"] = str(user_data["vlan_id"])
25
26     return 2 # OK
27
28 def match_policy(username, nas_ip, p):
29     # Query policy DB table, match conditions
30     # Return best-matching policy dict
31     return {
32         "auth_backend": "AD",
33         "check_blacklist": True
34     }
35
36 def get_backend(backend_type):
37     if backend_type == "AD":
38         return ADBackend()
39     elif backend_type == "TLS":
40         return TLSBackend()
41     elif backend_type == "ORACLE":
42         return OracleBackend()
43     else:
44         raise Exception("Unknown backend")
45
46 class ADBackend:
47     def authenticate(self, username, password):
48         # Do LDAP bind etc.
49         return True, {"vlan_id": 100, "mac": "00:11:22:33:44:55"}
50
51 # Add TLSBackend, OracleBackend, etc.
52

```

## 5. Call It in sites-enabled/default [↗](#)

Inside the `authorize` section:

```

1 authorize {
2     ...
3     python
4     ...
5 }

```

You can also use it in:

- `authenticate`

- `post-auth`
- `accounting`

## Advanced Enhancements [↗](#)

### Blacklist Grouping [↗](#)

Make the blacklist a table with:

- `id`
- `mac`, `username`, etc.
- `group_id`
- `reason`
- `expires_at`

Each policy can define:

- `blacklist_scope = "group" or "global"`

Python logic filters based on this.

### Failover Backends [↗](#)

If one backend (e.g., AD) fails, allow fallback to a secondary (e.g., local DB), configured per policy. Easy to express in a `backends: [AD, LOCAL]` list and iterate.

### Logging & Auditing [↗](#)

Log every auth attempt with:

- Matched policy ID
- User info
- Result (success/fail)
- Reason (blacklist, bad password, etc.)

Use syslog, flat file, or DB insert.

### Python Implementation (Dynamic Policy Matching) [↗](#)

```

1 import psycopg2
2
3 def match_policy(username, nas_ip, request_attrs):
4     conn = psycopg2.connect(...) # Your DB connection
5     cur = conn.cursor()
6
7     # Get all services
8     cur.execute("SELECT id, name, auth_backend, check_blacklist FROM rad_services")
9     services = cur.fetchall()
10
11     for service in services:
12         service_id, name, backend, check_bl = service
13         if service_matches(service_id, request_attrs, cur):
14             return {
15                 "id": service_id,
16                 "name": name,
17                 "auth_backend": backend,
```

```

18         "check_blacklist": check_bl
19     }
20
21     return None
22
23 def service_matches(service_id, attrs, cur):
24     cur.execute("""
25         SELECT attr_name, operator, attr_value
26         FROM rad_services_control
27         WHERE rad_service_id = %s
28     """, (service_id,))
29
30     conditions = cur.fetchall()
31
32     for attr_name, op, val in conditions:
33         req_val = attrs.get(attr_name)
34         if req_val is None:
35             return False
36
37         # Convert both to strings for simplicity
38         req_val = str(req_val)
39         val = str(val)
40
41         if not eval_condition(req_val, op, val):
42             return False
43
44     return True
45
46 def eval_condition(value, operator, expected):
47     if operator == "==":
48         return value == expected
49     elif operator == "!=":
50         return value != expected
51     elif operator == "startswith":
52         return value.startswith(expected)
53     elif operator == "endswith":
54         return value.endswith(expected)
55     elif operator == "contains":
56         return expected in value
57     elif operator == "in":
58         # expected = "val1,val2,val3"
59         return value in expected.split(",")
60     else:
61         return False

```

### Integration With FreeRADIUS Python Hook [🔗](#)

```

1 def authorize(p):
2     username = p.get("User-Name")
3     nas_ip = p.get("NAS-IP-Address")
4
5     # Convert request to flat dict of radius attributes
6     request_attrs = {k: str(v) for k, v in p.items()}
7
8     policy = match_policy(username, nas_ip, request_attrs)
9     if not policy:
10         return 1 # Reject
11

```

```

12     backend = get_backend(policy["auth_backend"])
13     success, user_data = backend.authenticate(username, p["User-Password"])
14
15     if not success:
16         return 1
17
18     if policy["check_blacklist"]:
19         if is_blacklisted(username, user_data["mac"]):
20             return 1
21
22     # Example reply
23     p["reply:Tunnel-Type"] = "VLAN"
24     p["reply:Tunnel-Medium-Type"] = "IEEE-802"
25     p["reply:Tunnel-Private-Group-Id"] = str(user_data["vlan_id"])
26
27     return 2

```

- For performance, you could **cache** `rad_services` and `rad_services_control` in memory or Redis and refresh every minute.

### Python Logic to Add Replies Dynamically [🔗](#)

Extend the `authorize()` function in the Python module:

```

1 def set_reply_attributes(service_id, p, cur):
2     cur.execute("""
3         SELECT attr_name, attr_value
4         FROM rad_service_replies
5         WHERE rad_service_id = %s
6     """, (service_id,))
7
8     for attr_name, attr_value in cur.fetchall():
9         # Explicitly set reply attributes
10        p[f"reply:{attr_name}"] = attr_value

```

### Full integration: [🔗](#)

```

1 def authorize(p):
2     username = p.get("User-Name")
3     nas_ip = p.get("NAS-IP-Address")
4
5     request_attrs = {k: str(v) for k, v in p.items()}
6
7     conn = psycopg2.connect(...) # reuse connection pool ideally
8     cur = conn.cursor()
9
10    policy = match_policy(username, nas_ip, request_attrs, cur)
11    if not policy:
12        return 1 # Reject
13
14    backend = get_backend(policy["auth_backend"])
15    success, user_data = backend.authenticate(username, p["User-Password"])
16
17    if not success:
18        return 1
19
20    if policy["check_blacklist"]:
21        if is_blacklisted(username, user_data["mac"]):

```

```
22         return 1
23
24     # Inject dynamic replies
25     set_reply_attributes(policy["id"], p, cur)
26
27     return 2 # OK
```

And use Python to evaluate conditions like:

```
1 def conditional_attr_match(attr_val, op, expected):
2     if op == "==":
3         return attr_val == expected
4     elif op == "startswith":
5         return attr_val.startswith(expected)
6     elif op == "!=":
7         return attr_val != expected
8     return False
```